

Conundrums in Unsupervised Keyphrase Extraction: Making Sense of the State-of-the-Art

Kazi Sandou Hasam and Vincent Ng

Human Language Technology Research Institute

University of Texas at Dallas

{sandou,vincent}@hlt.utdallas.edu

Abstract

State-of-the-art approaches for unsupervised keyphrase extraction are typically evaluated on a single dataset with a single parameter setting. Consequently, it is unclear how effective these approaches are on a new dataset from a different domain, and how sensitive they are to changes in parameter settings. To gain a better understanding of state-of-the-art unsupervised keyphrase extraction algorithms, we conducted a systematic evaluation and analysis of these algorithms on a variety of standard evaluation datasets.

1 Introduction

The keyphrases from a given document refer to a group of phrases that represent the document. Although we often come across texts from different domains such as scientific papers, news articles and blogs which are labeled with keyphrases by the authors, a large portion of the web content remains untagged. While keyphrases are excellent means for providing a concise summary of a document, recent research results have suggested that the task of automatically identifying keyphrases from a document is by no means trivial. Researchers have explored both supervised and unsupervised techniques to address the problem of automatic keyphrase extraction. Supervised methods typically cast this problem as a binary classification task, where a model is trained on annotated data to determine whether a given phrase is a keyphrase or not (e.g., Frank et al. (1999), Turney (2000, 2003), Huijin (2003), Medelyan et al. (2009)). A disadvantage of supervised approaches

is that they require a lot of training data and yet show bias towards the domain on which they are trained, undermining their ability to generalize well to new domains. Unsupervised approaches could be a viable alternative in this regard.

The unsupervised approaches for keyphrase extraction proposed so far have involved a number of techniques, including language modeling (e.g., Tomokiyama and Finin (2003)), graph-based ranking (e.g., Zhai (2002), Mihalcea and Tarau (2004), Wan et al. (2007)), twin and Xiao (2008), Jin et al. (2009a)), and clustering (e.g., Matsuo and Ishizuka (2004), Jin et al. (2009b)). While these methods have been shown to work well on a particular domain of text such as short paper abstracts and news articles, their effectiveness and portability across different domains have remained unexplored issues. Worse still, each of them is based on a set of assumptions, which may only hold for the dataset on which they are evaluated.

Consequently, we have little understanding of how effective the state-of-the-art systems would be on a completely new dataset from a different domain. A few questions arise naturally. How would these systems perform on a different dataset with their original configuration? What could be the underlying reasons in case they perform poorly? Is there any system that can generalize fairly well across various domains?

We seek to gain a better understanding of the state-of-the-art in unsupervised keyphrase extraction by examining the aforementioned questions. More specifically, we compare five unsupervised keyphrase extraction algorithms on four corpora with varying domains and statistical characteristics. These algorithms represent the im-

ion directions in this research area, including three and four recently proposed systems, namely, TextRank (Mihalcea and Tarau, 2004), SingleRank (Wan and Xiao, 2008), ExpandRank (Wan and Xiao, 2008), and a clustering-based approach (Liu et al., 2009b). Since none of these systems (except TextRank) are publicly available, we re-implement all of them and make them freely available for research purposes⁴. To our knowledge, this is the first attempt to compare the performance of state-of-the-art unsupervised keyphrase extraction systems on multiple datasets.

2 Corpora

Our four evaluation corpora belong to different domains with varying document properties. Table I provides an overview of each corpus.

The DUC2001 dataset (Over, 2001), which is a collection of 308 news articles, is annotated by Wan and Xiao (2008). We report results on all 308 articles in our evaluation.

The TIPSE dataset is a collection of 2,000 abstracts from journal papers including the paper titles. Each document has two sets of keyphrases assigned by the indexers: the *controlled keyphrases*, which are keyphrases that appear in the *abstracts*; and the *uncontrolled keyphrases*, which do not necessarily appear in the *abstracts*. This is a relatively popular dataset for automatic keyphrase extraction, as it was first used by Hotho (2003) and later by Mihalcea and Tarau (2004) and Liu et al. (2009b). In our evaluation, we use the set of 500 abstracts designated by these previous approaches as the test set and its set of uncontrolled keyphrases. Note that the average document length for this dataset is the smallest among all our datasets.

The NUS Keyphrase Corpus (Nguyen and Kan, 2007) includes 211 scientific conference papers with lengths between 4 to 12 pages. Each paper has one or more sets of keyphrases assigned by its authors and other annotators. We use all the 211 papers in our evaluation. Since the number of annotations can be different for different documents and the annotations are not specified along with the annotations, we decide to take the union

⁴See <http://www.htm.csail.mit.edu/nlp/corpora/>

of all the gold standard keyphrases from all the sets to construct one single set of annotation for each paper. As Table I shows, each NUS paper, both in terms of the average number of tokens (829) and candidate phrases (2027) per paper, is more than five times larger than any document from any other corpus. Hence, the number of candidate keyphrases that can be extracted is potentially large, making this corpus the most challenging of the four.

Finally, the ICSI meeting corpus (Liu et al., 2009a), which is annotated by Liu et al. (2009a), includes 161 meeting transcriptions. Following Liu et al., we remove topic segments marked as *topicchart* and *digraph* from the dataset and use all the remaining segments for evaluation. Each transcript contains three sets of keyphrases produced by the same three human annotators. Since it is possible to associate each set of keyphrases with its annotator, we evaluate each system on this dataset three times (once for each annotator) and report the average F-score. Unlike the other three datasets, the gold standard keys for the ICSI corpus are mostly unigrams.

3 Unsupervised Keyphrase Extractors

A generic unsupervised keyphrase extraction system typically operates in three steps (Section 3.1), which will help understand the unsupervised system explained in Section 3.2.

3.1 Generic Keyphrase Extractor

Step 1: Candidate lexical unit selection The first step is to filter out unnecessary words (tokens) from the input document and generate a list of potential keywords using heuristics. Commonly used heuristics include (1) using a stop word list to remove non-keywords (Nguyen et al., 2009b); and (2) allowing words with certain parts-of-speech tags (*DT*, *NN*, *ADJ*, *RB*, *VB*) to be considered candidate keywords (Mihalcea and Tarau, 2004; Liu et al., 2009a). Wan and Xiao (2008). In all of our experiments, we follow Wan and Xiao (2008) and select as candidate words with the following Penn Treebank tags (*NN*, *NN\$*, *NNP*, *NNPS*), and (3) which are obtained using the Stanford POS tagger (Bouhanova and Manning, 2000).

	Corpora			
Type	DUC2000	MSRA2007	NTUS	ICSI
#Documents	308	500	200	160
#Tokens/Document	876	1038	8290	1600
#Candidate words/Document	312	57	3270	463
#Candidate phrases/Document	207	194	2027	1296
#Tokens/Candidate phrase	1151	107	116	1151
#Gold keyphrases	2489	4916	2327	587
#Gold keyphrases/Document	8	9.8	11.0	3.6
LB/UB distribution (%)	11.6/11.8/19	11.6/12.6/19	11.6/16.6/17	6.8/19.1/11
#Tokens/Gold keyphrases	120	1213	120	113

Table 1: Corpus statistics for the four datasets used in this paper. A candidate word is typically a sequence of one or more adjectives and nouns. It is extracted from the document initially and considered a potential keyphrase. The LB/UB distribution indicates how the gold standard keys are distributed among unigrams, bigrams, trigrams, and other higher order n-grams.

Step 2: Lexical unit ranking Once the candidate list is generated, the next task is to rank these lexical units. To accomplish this, it is necessary to build a representation of the input text for the ranking algorithm. Depending on the underlying approach, each candidate word is represented by its syntactic and/or semantic relationship with other candidate words. The relationship can be defined using co-occurrence statistics, external resources (e.g., neighborhood, documents, Wikipedia), or other syntactic clues.

Step 3: Keyphrase formation In the final step, the ranked list of candidate words is used to form keyphrases. A candidate phrase, typically a sequence of nouns and adjectives, is selected as a keyphrase if (1) it includes one or more of the top-ranked candidate words (Vitaha and Tharu (2004)), (2) its rank (2009b)) or (2) the sum of the ranking scores of its constituent words makes it a top scoring phrase (Wan and Xiang (2008)).

3.2 The Five Keyphrase Extractors

As mentioned above, we re-implement five unsupervised approaches for keyphrase extraction. Below we provide a brief overview of each system.

3.2.1 Tf-Idf

Tf-Idf assigns a score to each term π in a document d , based on its frequency in d (term frequency) and how many other documents include π (inverse document frequency) and is defined as:

$$\text{tf-idf}_{\pi,d} = \text{tf}_{\pi,d} \times \log(D/d_f) \quad (1)$$

where D is the total number of documents and d_f is the number of documents containing π .

Given a document, we first compute the Tf-Idf score of each candidate word (see Step 1 of the generic algorithm). Then, we extract all the longest n-grams consisting of candidate words and score each n-gram by summing the Tf-Idf scores of its constituent unigrams. Finally, we output the top N n-grams as keyphrases.

3.2.2 TextRank

In the TextRank algorithm (Mihalcea and Tarau (2004)), a text is represented by a graph. Each vertex corresponds to a word type. A weight w_{ij} is assigned to the edge connecting the two vertices v_i and v_j and its value is the number of times the corresponding word types co-occur within a window of m words in the associated text. The goal is to (1) compute the score of each vertex which reflects its importance and then (2) use the word types that correspond to the highest scored vertices to form keyphrases for the text. The score for v_i , $S(v_i)$, is initialized with a default value and is computed in an iterative manner until convergence using this recursive formula:

$$S(v_i) = (1 - d) + d \times \sum_{v_j \in \text{Ad}(v_i)} \frac{w_{ji}}{\sum_{v_k \in \text{Ad}(v_j)} w_{kj}} S(v_j) \quad (2)$$

where $\text{Ad}(v_i)$ denotes v_i 's neighbors and d is the damping factor set to 0.85 (Brin and Page (1998)). Intuitively, a vertex will receive a high score if it has many high-scored neighbors. As noted before, after convergence, the 7% top-scored vertices are

selected as keywords. Adjacent keywords are then collapsed and output as a keyphrase.

According to Viñalsica and Torrau (2004), TextRank is the score on the Wikipedia dataset is achieved when only nouns and adjectives are used to create a uniformly weighted graph for the text under consideration, where an edge connects two word types only if they co-occur within a window of two words. Hence, our implementation of TextRank follows this configuration.

3.2.3 SingleRank

SingleRank (Wan and Xiao, 2008) is essentially a TextRank approach with three major differences. First, while each edge in a TextRank graph (in Viñalsica and Torrau's implementation) has the same weight, each edge in a SingleRank graph has a weight equal to the number of times the two corresponding word types co-occur. Second, while in TextRank only the word types that correspond to the top-ranked vertices can be used to form keyphrases, in SingleRank, we do not filter out any low-scored vertices. Rather, we (1) score each candidate keyphrase, which can be any longest-matching sequence of nouns and adjectives in the text under consideration, by summing the scores of its constituent word types obtained from the SingleRank graph and (2) output the N highest-scored candidates as the keyphrases for the text. Finally, SingleRank employs a window size of 10 rather than 2.

3.2.4 ExpandRank

ExpandRank (Wan and Xiao, 2008) is a TextRank extension that exploits neighbourhood knowledge for keyphrase extraction. For a given document d_i , the approach first finds its k nearest neighbouring documents from the accompanying document collection using a similarity measure (e.g., cosine similarity). Then, the graph for this is built using the co-occurrence statistics of the candidate words collected from the document in itself and its k nearest neighbors.

Specifically, each document is represented by a term vector where each vector dimension corresponds to a word type present in the document and its value is the Tfidf score of that word type for that document. For a given document d_i , its k

nearest neighbors are identified and together they form a larger document set of $k+1$ documents, $D = \{d_1, d_2, \dots, d_k\}$. Given this document set, a graph is constructed, where each vertex corresponds to a candidate word type in D , and each edge connects two vertices v_i and v_j if the corresponding word types co-occur within a window of m words in the document set. The weight of an edge $w(v_i, v_j)$ is computed as follows:

$$w(v_i, v_j) = \sum_{d \in D} \text{sim}(v_i, v_j) \times \text{freq}_{v_i, v_j}(d, D) \quad (6)$$

where $\text{sim}(v_i, v_j)$ is the cosine similarity between v_i and v_j and $\text{freq}_{v_i, v_j}(d, D)$ is the co-occurrence frequency of v_i and v_j in document d . Once the graph is constructed, the rest of the procedure is identical to SingleRank.

3.2.5 Clustering-based Approach

Ji et al. (2009) propose to cluster candidate words based on their semantic relationship to ensure that the extracted keyphrases cover the entire document. The objective is to have each cluster represent a unique aspect of the document and take a representative word from each cluster so that the document is covered from all aspects.

More specifically, their algorithm (henceforth referred to as KeyCluster) first filters out the stop words from a given document and treats the remaining n-grams as candidate words. Second, for each candidate, its relatedness with another candidate is computed by (1) counting how many times they co-occur within a window of size m in the document and (2) using Wikipedia-based statistics. Third, candidate words are clustered based on their relatedness with other candidates. Three clustering algorithms are used of which spectral clustering yields the best score. Once the clusters are formed, one representative word, called an exemplar term, is picked from each cluster. Finally, KeyCluster extracts from the document all the longest n-grams starting with zero or more adjectives and ending with one or more nouns, and if such an n-gram includes one or more exemplar words, it is selected as a keyphrase. As a post-processing step, a frequent word list generated from Wikipedia is used to filter out the frequent n-grams that are selected as keyphrases.

4 Evaluation

4.1 Experimental Setup

TextRank and SingleRank setup. Following Mihalcea and Tarau (2004) and Wan and Xiao (2008), we set the co-occurrence window size for TextRank and SingleRank to 2 and 10, respectively, as these parameter values have yielded the best results for the evaluation datasets.

ExpandRank setup. Following Wan and Xiao (2008), we find the 5 nearest neighbors for each document from the remaining documents in the same corpus. The other parameters are set in the same way as in SingleRank.

KeyCluster setup. As argued by Jin et al. (2009b), Wikipedia-based relatedness is computationally expensive to compute. As a result, we follow them by computing the n-gram-based relatedness instead, using a window of size 10. Then, we cluster the candidate words using spectral clustering, and use the frequent word lists that they generously provided us to post-process the resulting keyphrases by filtering out those that are frequent unigrams.

4.2 Results and Discussion

In an attempt to gain a better insight into the five unsupervised systems, we report their performance in terms of precision/recall curves for each of the four datasets (see Figure 1). This contrasts with essentially all previous work, where the performance of a keyphrase extraction system is reported in terms of an F-score obtained via a particular parameter setting on a particular dataset. We generate the curves for each system as follows. For TextRank, SingleRank, and ExpandRank, we vary the number of keyphrases (N) predicted by each system. For TextRank, instead of varying the number of predicted keyphrases, we vary π , the percentage of top-scored vertices (i.e., unigrams) that are selected as keywords at the end of the ranking step. The reason is that TextRank only imposes a ranking on the unigrams but not on the keyphrases generated from the high-ranked unigrams. For KeyCluster, we vary the number of clusters produced by spectral clustering rather than the number of predicted keyphrases, again because KeyCluster does not impose a ranking on

the resulting keyphrases. In addition to give an estimate of how each system performs in terms of F-score, we also plot curves corresponding to different F-scores in these graphs.

TextRank. Consistent with our intuition, the precision of TextRank drops as recall increases. Although this is the simplest of the five approaches, TextRank is the best performing system on all but the *Inspe* dataset, where TextRank and KeyCluster beat it in just a few cases. It clearly outperforms all other systems for NUS and TCM.

TextRank. The TextRank curves show a different progression than TextRank: precision does not drop as much when recall increases. For instance, in case of PUC and ICSI, precision is not sensitive to changes in recall. Perhaps somewhat surprisingly, its precision increases with recall for *Inspe*, allowing it to even reach a point (towards the end of the curve) where it beats TextRank. While additional experiments are needed to determine the reason for this somewhat counter-intuitive result, we speculate that this may be attributed to the fact that the TextRank curves are generated by progressively increasing π (i.e., the percentage of top-ranked vertices/unigrams that are used to generate keyphrases) rather than the number of predicted keyphrases, as mentioned before. Increasing π does not necessarily imply an increase in the number of predicted keyphrases, however. To see the reason, consider an example in which we want TextRank to extract the keyphrase “advanced machine learning” from a given document. Assume that TextRank ranks the unigrams “advanced”, “learning”, and “machine” first, second, and third, respectively, in its ranking step. When $\pi = \frac{2}{3}$, where n denotes the total number of candidate unigrams, only the two highest-ranked unigrams (i.e., “advanced” and “learning”) can be used to form keyphrases. This implies that “advanced” and “learning” will each be predicted as a keyphrase, but “advanced machine learning” will not. However, when $\pi = \frac{3}{3}$, all three unigrams can be used to form a keyphrase, and since TextRank collapses unigrams adjacent to each other in the text to form a keyphrase, it will correctly predict “advanced machine learning” as a keyphrase. Note that as we increase π from $\frac{2}{3}$ to $\frac{3}{3}$ (recall increases), and so does precision, since

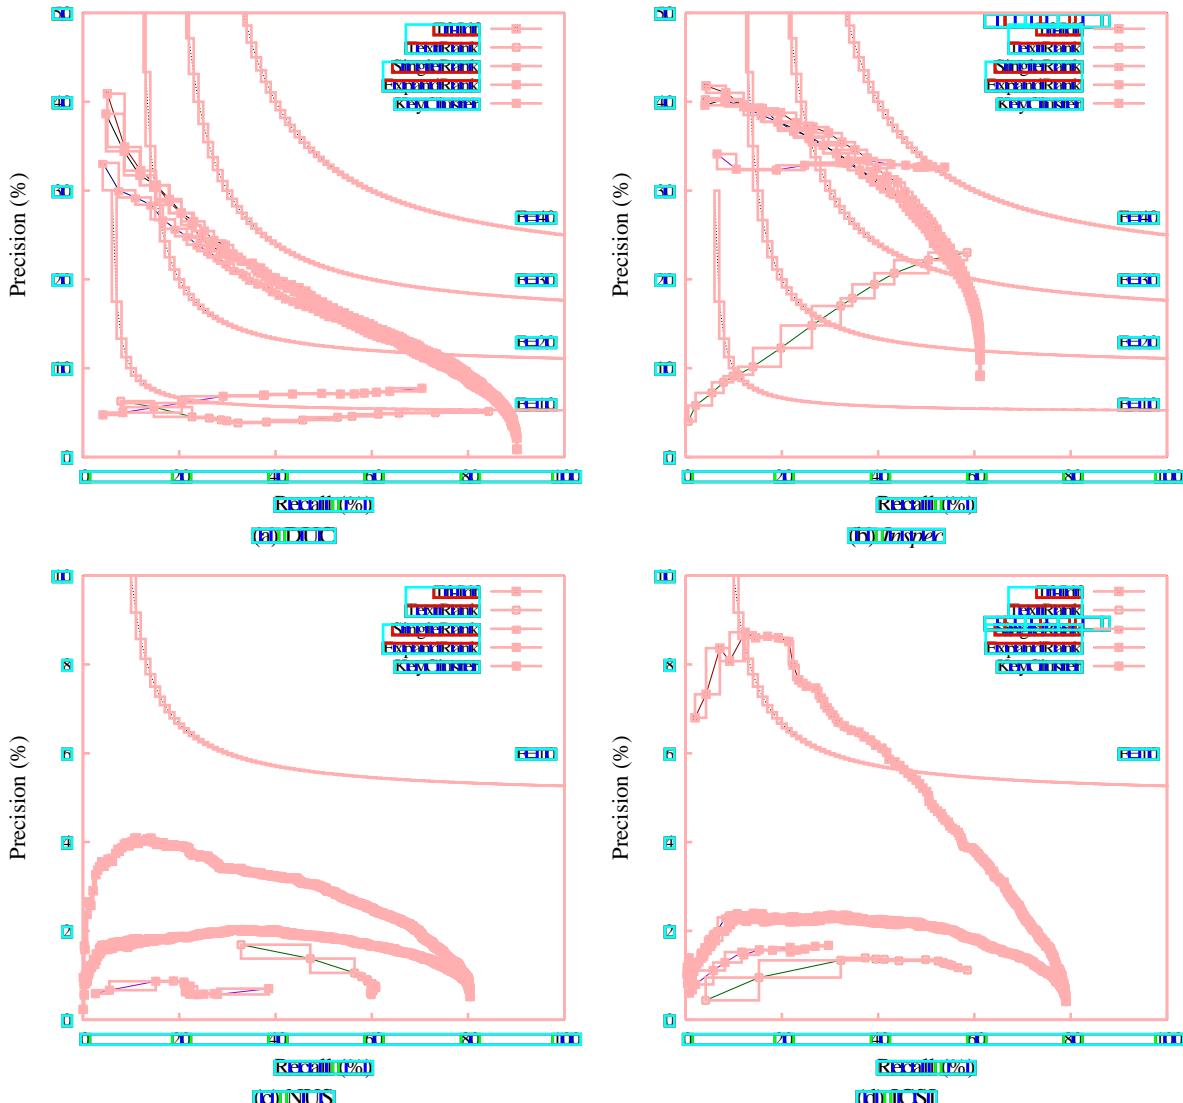


Figure 11: Precision-recall curves for all four datasets

'advanced' and 'learning' are now combined to form one keyphrase (and hence the number of predicted keyphrases decreases). In other words, it is possible to see a simultaneous rise in precision and recall in a TextRank curve. A natural question is: why does it happen only for *Inspec* but not the other datasets? The reason could be attributed to the fact that *Inspec* is composed of abstracts, since the number of keyphrases that can be generated from these short documents is relatively small; precision may not drop as severely as with the other datasets even when all of the n-grams are used to form keyphrases.

On average, TextRank performs much worse

compared to TriRank. The curves also prove TextRank's sensitivity to *Thom Yorke*, but not on the other datasets. This certainly gives more insight into TextRank since it was evaluated on *Inspec* only (to 73%) by Mihalcea and Tarau (2004).

SingleRank—SingleRank, which is supposed to be a simple variant of TextRank, surprisingly exhibits very different performance. First, it shows a more intuitive initial precision drop as recall increases. Second, SingleRank outperforms TextRank by big margins on all the datasets. Later we will examine which of the differences between them is responsible for the differing performance.

	DUC	NYUW	NUS	CSC		
Parameter	R	Parameter	R	Parameter	R	
Mihalcea	1/100	1/100	1/100	1/100	1/100	
TextRank	1/1000%	1/100	1/1000%	1/100	1/1000%	1/100
SingleRank	1/100	1/100	1/100	1/100	1/100	
ExpandRank	1/100	1/100	1/100	1/100	1/100	
KeyCluster	1/100000	1/100	1/100000	1/100	1/100000	1/100

Table 2: Best parameter settings. R is the number of predicted keyphrases, 1/T is the percentage of vertices selected as keywords in TextRank, m is the number of clusters in KeyCluster expressed in terms of the fraction of candidate words.

ExpandRank: Consistent with Wan and Xiao (2008), ExpandRank beats SingleRank on DUC when a small number of phrases are predicted, until their difference diminishes as more phrases are predicted. On the other hand, their performance is indistinguishable from each other on the other three datasets. A natural question is: why does ExpandRank improve over SingleRank only for DUC but not for the other datasets? To answer this question, we look at the DUC articles and find that in many cases, the nearest neighbors of a document are on the same topic involving the same entities as the document itself, presumably because many of these news articles are simply updated versions of an evolving event. Consequently, the graph built from the neighboring documents is helpful for predicting the keyphrases of the given document. Such topic-wise similarity among the nearest documents does not exist in the other datasets, however.

KeyCluster: As in TextRank, KeyCluster does not always yield a drop in precision as recall improves. This again may be attributed to the fact that the KeyCluster curves are generated by varying the number of clusters rather than the number of predicted keyphrases, as well as the way keyphrases are formed from the exemplars. Another reason is that the frequent Wikipedia n-grams are excluded during post-processing, making KeyCluster more resistant to precision drops. Overall, KeyCluster performs slightly better than TextRank on DUC and CSC, yields the worst performance on NUS, and scores the best on NYUW when the number of clusters is high. These results seem to suggest that KeyCluster works better if more clusters are used.

Best parameter settings: Table 2 shows for each system the parameter values yielding the best F1 score on each dataset. Two points deserve men-

tion: First, in comparison to SingleRank and ExpandRank, Mihalcea outputs fewer keyphrases to achieve its best F1 score on most datasets. Second, the systems output more keyphrases on NUS than on other datasets to achieve their best F1 scores (e.g., 60 for Mihalcea, 190 for SingleRank, and 1177 for ExpandRank). This can be attributed in part to the fact that the F1 scores on NUS are low for all the systems and exhibit only slight changes as we output more phrases.

Our reimplementations: Do our duplicated systems yield scores that match the original scores? Table 3 sheds light on this question.

First, consider KeyCluster, where our scores lag behind the original score by approximately 5%. An examination of Lin et al.'s (2009b) results reveals a subtle caveat in keyphrase extraction evaluations. Lin et al. not all gold-standard keyphrases appear in their associated document and as a result, none of the five systems we consider in this paper can achieve a recall of 100% while Mihalcea and Tarau (2004) and our reimplementations use *all* of these gold-standard keyphrases in our evaluation. Huihui (2003) and Lin et al. address this issue by using as gold-standard keyphrases only those that appear in the corresponding document when computing recall. This explains why our KeyCluster score (38.9%) is lower than the original score (43.9%). If we follow Lin et al.'s way of computing recall, our reimplementation score goes up to 42.4%, which lags behind their score by only 1.1%.

Next, consider TextRank, where our scores lag behind Mihalcea and Tarau's original score by more than 25 points. We verified our implementation against a publicly available implementation

that is based on Mihalcea and Tarau's scores are not directly comparable but in general did not point this out while computing scores in their paper.

Dataset	F-score	
	Original	Ours
Wicihi	DUO	26.41
TextRank	Imp2pt	36.02
SingleRank	DUO	27.02
ExpandRank	DUO	36.07
KeyCluster	Imp2pt	48.96
		38.19

Table 3: Original vs our implementation scores

of TextRank², and are confident that our implementation is correct. It is also worth mentioning that using our reimplementation of SingleRank we are able to match the best scores reported by Wu and Xie (2008) and Wu et al. (2009) on *Inspec*.

We score 2 and 5 points less than Wu and Xie (2008) implementations of SingleRank and ExpandRank respectively. We speculate that document pre-processing (e.g., stemming) has contributed to the discrepancy. (Our additional experiments are needed to determine the reason.)

SingleRank vs TextRank. Figure 1 shows that SingleRank behaves very differently from TextRank. As mentioned in Section 3.1/2/3, the two algorithms differ in three major aspects. To determine which aspect is chiefly responsible for the large difference in their performance, we conduct three ‘ablation’ experiments. Each experiment modifies exactly one of these aspects in SingleRank so that it behaves like TextRank, effectively ensuring that the two algorithms differ only in the remaining two aspects. More specifically, in the three experiments, we (i) change SingleRank’s window size to 1; (ii) build an unweighted graph for SingleRank; and (iii) incorporate TextRank’s way of forming keyphrases into SingleRank, respectively. Figure 1 shows the resultant curves along with the SingleRank and TextRank curves on *Inspec* taken from Figure 1b. As we can see, the way of forming phrases, rather than the window size or the weight assignment method, has the largest impact on the scores. In fact, after incorporating TextRank’s way of forming phrases, SingleRank exhibits a remarkable drop in performance, yielding a curve that resembles the TextRank curve. Also note that SingleRank achieves better recall values than TextRank. To see the reason, recall that TextRank requires that every word of a gold keyphrase must appear among the top-

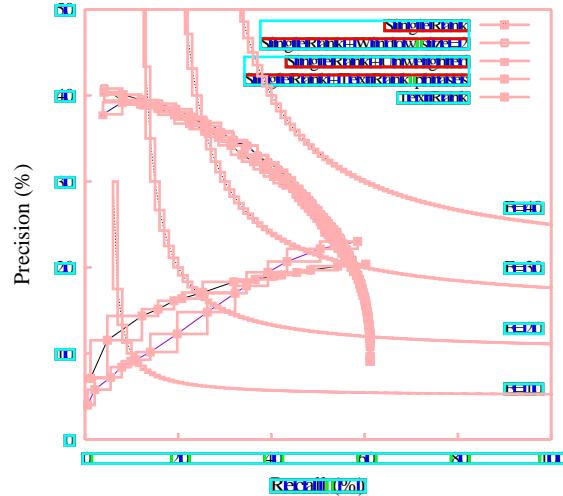


Figure 2: Ablation results for SingleRank on *Inspec*

ranked unigrams. This is a fairly strict criterion, especially in comparison to SingleRank, which does not require all n-grams of a gold keyphrase to be present in the top-ranked list. (We observe similar trends for the other datasets.)

5 Conclusions

We have conducted a systematic evaluation of five state-of-the-art unsupervised keyphrase extraction algorithms on datasets from four different domains. Several conclusions can be drawn from our experimental results. First, to fully understand the strengths and weaknesses of a keyphrase extractor, it is essential to evaluate it on multiple datasets. In particular, evaluating it on a single dataset has proven inadequate, as good performance can sometimes be achieved due to certain statistical characteristics of a dataset. Second, as demonstrated by our experiments with TextRank and SingleRank, post-processing steps such as the way of forming keyphrases can have a large impact on the performance of a keyphrase extractor. Hence, it may be worthwhile to investigate alternative methods for extracting candidate keyphrases (e.g., Kumar and Srinivasan (2008); You et al. (2009)). Finally, despite the large amount of recent work on unsupervised keyphrase extraction, our results indicated that Tf-IDf remains a strong baseline, offering very robust performance across different datasets.

²<http://github.com/ishareefhisi/TextRank>

Acknowledgments

We thank the three anonymous reviewers for their comments! Many thanks to Annette Huhin and Ying Liu for providing us with the *VySpEd* and ICSI datasets; Radu Mihalcea, Parag Nathan and Xiaojun Wan for helping us understand their algorithms/implementation; and Peng Ji for providing us with the frequent word lists that he and his co-authors used in their paper! This work was supported in part by NSF Grant IIS-0812261.

References

- Brun Sergey and Lawrence Page. 1998. The anatomy of a large-scale hypertextual Web search engine. *Computer Networks* 30(1-7) (1998), 107-117.
- Franki, Elke, Gordon Wilfayken, Ian H. Witten, Carl Gutwin, and Craig G. Nevill-Manning. 1999. Domain-specific keyphrase extraction. In *Proceedings of the 16th International Joint Conference on Artificial Intelligence*, pages 668-673.
- Huhin, Anette. 2003. Improved automatic keyword extraction given more linguistic knowledge. In *Proceedings of the 2003 Conference on Empirical Methods in Natural Language Processing*, pages 216-220.
- Jainini, Adam, Don Barou, Jane Edwards, Dan Flesch, David Gelbart, Nelson Moreira, Barbara Peleskin, Thilo Pfandl, Elizabeth Shriberg, Andreas Stolcke, and Chuck Wooters. 2003. The ICSI meeting corpus. In *Proceedings of 2003 IEEE Conference on Automatic Speech and Signal Processing*, pages 364-367.
- Kumar, Nitaj and Kannan Srinivasan. 2008. Automatic keyphrase extraction from scientific documents using negation filtration techniques. In *Proceedings of the Eighth ACM Symposium on Document Engineering*, pages 199-208.
- Jin, Feilai, Deana Pennell, Fei Liu, and Ying Liu. 2009a. Unsupervised approaches for automatic keyword extraction using meeting transcripts. In *Proceedings of Human Language Technologies: The 2009 Annual Conference of the North American Chapter of the Association for Computational Linguistics*, pages 620-628.
- Jin, Zhiyan, Peng Jin, Yabin Zheng, and Mansong Sun. 2009b. Clustering to find exemplar terms for keyphrase extraction. In *Proceedings of the 2009 Conference on Empirical Methods in Natural Language Processing*, pages 757-766.
- Mansouf, Yutaka and Mutsuo Ishizuka. 2004. Keyword extraction from a single document using word co-occurrence statistical information. *International Journal on Artificial Intelligence Tools* 13(01):1-169.
- Meredith, Oren, Elke Franki, and Ian H. Witten. 2009. Human-computing tagging using automatic keyphrase extraction. In *Proceedings of the 2009 Conference on Empirical Methods in Natural Language Processing*, pages 1318-1327.
- Mihalcea, Rada and Pantel P. Mihalcea. 2004. Textrank: Bringing order into texts. In *Proceedings of the 2004 Conference on Empirical Methods in Natural Language Processing*, pages 404-410.
- Nguyen, Thuy Dung and Minh-Hien Kan. 2007. Keyphrase extraction in scientific publications. In *Proceedings of the International Conference on Asian Digital Libraries*, pages 3117-3126.
- Oveli, Panit. 2001. Introduction to DUC2001: An intrinsic evaluation of generic news text summarization systems. In *Proceedings of the 2001 Document Understanding Conference*.
- Tomokiyama, Takashi and Matthew Hurst. 2003. A language model approach to keyphrase extraction. In *Proceedings of the 2003 Workshop on Information Extraction*.
- Tourneau, Kristina and Christopher D. Manning. 2000. Enriching the knowledge sources used in a maximum entropy part-of-speech tagger. In *Proceedings of the 2000 Joint SIGART Conference on Empirical Methods in Natural Language Processing and Very Large Corpora*, pages 63-70.
- Turvey, Peter. 2000. Learning algorithms for keyphrase extraction. *Information Retrieval* 2(3):3-330.
- Turvey, Peter. 2003. Coherent keyphrase extraction via web mining. In *Proceedings of the 18th International Joint Conference on Artificial Intelligence*, pages 434-439.
- Wan, Xiaojun and Liangmo Xie. 2008. Single document keyphrase extraction using neighborhood knowledge. In *Proceedings of the 2008 AAAI Conference on Artificial Intelligence*, pages 855-860.
- Wan, Xiaojun, Liangmo Xie, and Liangmo Xie. 2007. Towards an iterative reinforcement approach to tie multiple document summarization and keyword extraction. In *Proceedings of the 24th Annual Meeting of the Association of Computational Linguistics*, pages 652-659.
- Yiou, Wei, Dominique Fontaine, and Jean-Paul Bapteste. 2009. Automatic keyphrase extraction with a refined candidate set. In *Proceedings of the 2009 IJCNLP-2009 International Joint Conference on Web Intelligence and Intelligent Agent Technologies*, pages 517-519.
- Zhai, Hongyuan. 2002. Generic summarization and keyphrase extraction using mutual reinforcement principle and sentence clustering. In *Proceedings of the 20th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 1113-1120.